




Multi-Layer Feed-Forward Neural Networks

The most common case involves a single hidden layer:




Multi-Layer Feed-Forward Neural Networks

Each *hidden unit* can be considered as single output perceptron network:




Multi-Layer Feed-Forward Neural Networks


The output unit of the multi-layer network can then be considered a soft-thresholded linear combination of the hidden units (which are equivalent to single output unit perceptrons):



Multi-Layer Feed-Forward Neural Networks



(a)



(b)

Figure 20.23 (a) The result of combining two opposite-facing soft threshold functions to produce a ridge. (b) The result of combining two ridges to produce a bump.

Expressibility of Multi-Layer Neural Networks

- With a single, sufficiently large hidden layer, it is possible to represent *any continuous* function of the inputs with *arbitrary* accuracy.
- Unfortunately, for any *particular* network, it is harder to characterize exactly which functions can be represented and which ones cannot.
- As a consequence, given a particular learning problem, it is unknown how to choose the *right number of hidden units* in advance.
- One usually resorts to cross validation, but this can be computationally expensive for large networks.

Back-Propagation Learning

- We need to consider multiple output units for multi-layer networks. Let (\mathbf{x}, \mathbf{y}) be a single sample with its desired output labels $\mathbf{y} = \{y_1, \dots, y_i, \dots, y_M\}$.
- The error at the output units is just $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$, and we can use this to adjust the weights between the hidden layer and the output layer.
- The above steps produces a term equivalent to the error at the hidden layer, i.e. the error at the output layer is **back-propagated** to the hidden later.
- This is subsequently used to update the weights between the input units and the hidden layer.

Back-Propagation Learning in Detail

Step 1: Update the weights between the hidden and output layers.

- Let Err_i be the i -th component of the error vector $\mathbf{y} - h_{\mathbf{W}}(\mathbf{x})$.
- Define $\Delta_i = Err_i \times g'(in_i)$.
- The weight-update rule becomes

$$W_{j,i} \longleftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

This is similar to weight-updates for Perceptrons!

Back-Propagation Learning in Detail

Step 2: Back-propagate the error to the hidden layer.

- The idea is that the hidden node j is “responsible” for some fraction of the error Δ_i in each of the output nodes to which it connects.
- Thus the Δ_i values are divided according to the strength (weight) of the connection between the hidden node and the output node:

$$\Delta_j = g'(in_j) \sum_i W_{j,i} \Delta_i$$

Back-Propagation Learning in Detail

Step 3: Update the weights between the input units and the hidden layer.

- Again, this is similar to weight-updates in Perceptrons:

$$W_{k,j} \longrightarrow W_{k,j} + \alpha \times a_k \times \Delta_j$$

Summary of Back-Propagation Learning

For the general case of *multiple hidden* layers:

- ① Compute the Δ values for the output units, using the observed error.
- ② Starting with the output layer, repeat the following for each layer in the network, until the earliest hidden layer is reached:
 - Propagate the Δ values back to the previous layer.
 - Update the weights between the two layers.
- ③ Repeat Steps 1 to 2 for all training samples.

Algorithm for Back-Propagation Learning

```

function BACK-PROP-LEARNING(examples, network) returns a neural network
  inputs: examples, a set of examples, each with input vector  $\mathbf{x}$  and output vector  $\mathbf{y}$ 
           network, a multilayer network with  $L$  layers, weights  $W_{j,i}$ , activation function  $g$ 

  repeat
    for each  $e$  in examples do
      for each node  $j$  in the input layer do  $a_j \leftarrow x_j[e]$ 
      for  $\ell = 2$  to  $M$  do
         $in_i \leftarrow \sum_j W_{j,i} a_j$ 
         $a_i \leftarrow g(in_i)$ 
      for each node  $i$  in the output layer do
         $\Delta_i \leftarrow g'(in_i) \times (y_i[e] - a_i)$ 
      for  $\ell = M - 1$  to  $1$  do
        for each node  $j$  in layer  $\ell$  do
           $\Delta_j \leftarrow g'(in_j) \sum_i W_{j,i} \Delta_i$ 
          for each node  $i$  in layer  $\ell + 1$  do
             $W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$ 
      until some stopping criterion is satisfied
  return NEURAL-NET-HYPOTHESIS(network)
  
```

Figure 20.25 The back-propagation algorithm for learning in multilayer networks.

Results of Back-Propagation Learning




Figure 20.26 (a) Training curve showing the gradual reduction in error as weights are modified over several epochs, for a given set of examples in the restaurant domain. (b) Comparative learning curves showing that decision-tree learning does slightly better than back-propagation in a multilayer network.